

Learning from Data – Week 2

Assignment 2: Decision Tree, Random Forest, K-NN

General remarks

This assignment is meant to get you further acquainted with scikit-learn and in particular the principles of the algorithms we have discussed. You will also explore a little further the contribution of features and varying parameters in running classification experiments. In the discussion of the results bear in mind the theoretical points we have touched on in class.

Specifically, for the practical parts, you are asked to train and test a few models using different algorithms on the same kind of data that we used for Assignment 1, but **only considering the six-way classification**. Exploiting coding you have done for Assignment 1 regarding measures, you will have to discuss how the different models perform, and why.

What you have to do:

- Upload on Nestor the code with your best model (see Exercise 2.3).
Important: You have to assume that we will run it using arguments like this:
`LFD_assignment2.py -i <trainset> -ts <testset>`
where trainset and testset have the same format. You are given the trainset, which is the same dataset as last week, but we're holding out the test set. The code itself will be a (small) part of the final grade, so make sure to structure it and add comments!
- Fill in a Google form where you will answer a few questions: <https://forms.gle/6dDg5yZJHDrZy5py6>. When filling this in, please refer to this pdf for more details on the questions, instructions and clarifications.

Deadline: 20 September, 10:59 AM.

This assignment is to be completed **individually**.

Data

For this assignment, we will be using the collection of reviews we used for Assignment 1. The labels we are interested in are the six denoting the topics: **books**, **camera**, **dvd**, **health**,

music, software. Likely you already downloaded the data, but it's also still available on Nestor (Week 1).

Exercise 2.1 – Decision Tree

2.1.1 Theory

| colour | size | shape | edible |
|--------|-------|-----------|--------|
| yellow | small | round | yes |
| yellow | small | round | no |
| green | small | irregular | yes |
| green | large | irregular | no |
| yellow | large | round | yes |
| yellow | small | round | yes |
| yellow | small | round | yes |
| yellow | small | round | yes |
| green | small | round | no |
| yellow | large | round | no |
| yellow | large | round | yes |
| yellow | large | round | no |
| yellow | large | round | no |
| yellow | large | round | no |
| yellow | small | irregular | yes |
| yellow | large | irregular | yes |

Consider the dataset on the left. Given that the last column is the class you want to predict, what is the tree that best represents this data? You have to figure it out without running it through scikit-learn! It's actually a good idea to draw it by hand, thinking what the first branch will be, and why. If you want to show calculations you can add them in the picture.

2.1.2 Practice

Scikit-learn provides an implementation of a decision tree:

```
from sklearn.tree import DecisionTreeClassifier
```

scikit-learn.org/stable/modules/tree.html

Pruning isn't supported in scikit-learn, but you can set several parameters of the algorithm, such as tree size, minimal number of samples for a split, and so on. You can check the documentation and slides for more info on this. Please, run the relevant experiments and answer the questions about pruning and about changing parameters in the Google form.

It's also possible to use the Random Forest algorithm in a similar way:

```
from sklearn.ensemble import RandomForestClassifier
```

<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>

There are no related questions in the Google form, but please implement this algorithm in the submitted code and play around a bit with the parameters. You have to use it in Exercise 2.3 anyway.

Exercise 2.2 – K-Nearest Neighbor

Scikit-learn also provides an implementation of the K-nearest neighbor algorithm:

```
from sklearn.neighbors import KNeighborsClassifier
```

scikit-learn.org/stable/modules/neighbors.html#nearest-neighbors-classification

Run a K-NN classifier on the data performing a six-way classification. Experiment with different values of K seeing what happens on your development set. Plot accuracy and f-scores on a graph, you can upload this on the Google form under question KNN-1. The following questions in the form also have to do with varying K and varying performance.

Exercise 2.3 – Comparisons

2.3.1 Running times

Check the time it takes to train and test the model, and compare it with what you observed for the Decision Tree, Random Forest (see above) and Naive Bayes. On the form, there is a question where you have to explain why the algorithms have different run times at train and test time.

A way to check time (but you can choose whatever you prefer):

```
import time

t0 = time.time()
classifier.fit(X, y)
train_time = time.time() - t0
print("training time: ", train_time)
```

You should do the same at test time, and for both models, to compare.

2.3.2 Best Model

This part of the assignment gives you some more freedom: please give it a shot to develop the best model possible given the training data! Note that it is now impossible to cheat: your model will be run on data that we held out on purpose, i.e. data you have not trained nor developed on. We will run this model for you, so make sure it works like this:

```
python LFD_assignment2.py -i <trainset> -ts <testset>
```

Now, how can you make a better model? Say that you're working with a decision tree or random forest, you can modify parameter values such as those mentioned above. If you're

working with K-NN, you can choose an optimal K. How do you *tune* such parameters? How do you choose the *optimal* K? By cross-validating your training set or testing on your development set! Also, you can still decide you'd rather use Naive Bayes, maybe with a selection of features (see below)?

We will go into this in more detail next week, but it's encouraged to get a little bit more sophisticated on the features you're using. The Assignment 1 script in its original form takes all words as features. Do you want to refine this? Some ideas to get you going:

- Look at the possible arguments of CountVectorizer and TfidfVectorizer and experiment with them
- Combine different vectorizers (hint: use FeatureUnion)
- Filtering out stopwords
- Lemmatizing or stemming the input words (NLTK)
- Using a tagger (Named Entity, POS-tags) to add as input features
- Combining the predictions of different models. You can of course combine the output of different algorithms, but also train multiple versions of the same algorithm with different parameter settings

An important challenge as well is to prevent overfitting! Take a look at the advice from week 1 and week 2 and make sure you have a way to prevent it as much as possible.

It's entirely up to you to choose which model you want us to run on test data, but: it has to be an algorithm we discussed in class already (so no SVM or neural networks).

Please, upload on Nestor the script that runs your best model, and provide a brief summary of your choices (algorithm, parameters, features) in the Google form, where asked.

You can check that everything works by splitting the data and saving it in different files (a train file and a test file). **Important:** we also expect your script to output the standard evaluation measures you have already used for Assignment 1. The students that scored best on these metrics get a bonus!